

Structural Estimation — Result Tables

23c783b-dirty 2026-08-06 13:57

3 tables

Contents

1	baseline_college_work_choice	2
2	baseline_outcomes	3
3	numerical_diagnostics	4

1 baseline_college_work_choice

Table 1: College vs. Work Decisions in the Baseline

Decision	Number	Share (%)
College	731	14.6
Work	4,269	85.4
Total	5,000	100.0

Note: Number and share of individuals choosing each post-family-stage path in the baseline simulation.

2 baseline_outcomes

Table 2: End-of-Family Outcomes: Baseline

Outcome (at $t = 18$)	Baseline
Mean final assets ^a	196.04
Mean final human capital	3.3074

^a Values in thousands of U.S. dollars.

Note: Means computed over the simulated cohort at the end of the family stage.

3 numerical_diagnostics

Table 3: Numerical Diagnostics

Diagnostic	Value
Minimum converged share (parent solve)	1.0000
Bellman residual, max (child work path)	5.73e-13
Bellman residual, mean	9.95e-14
Monotonicity violations, V in assets	0.03%
Monotonicity violations, c in assets	3.72%
Simulated assets below $a_{\min}$	0.00%
Simulated assets above $a_{\max}$	0.00%
Simulated states non-finite	0.00%
Stored transitions off-grid, assets	1.07% (max 1.0e-06)
Stored transitions off-grid, HC	2.00% (max 1.0e-03)
$k_{\max}$ ceiling binds	0.00%
Bellman optimality residual, max	1.50e-05
States improved by re-optimization	2.00%
Linear vs. cubic continuation, max $ \Delta c $	5.46e+00
Linear vs. cubic continuation, max $ \Delta h $	1.08e-01
Mean terminal parental assets	19.6040
Bootstrap standard error	0.1335

Note: Computed on the solved model and the simulated cohort. Bellman residuals are relative; standard errors are bootstrapped over 200 resamples. The optimality residual re-optimizes sampled states from four starts and compares the maximum against the stored value, so unlike the consistency residual it detects a suboptimal policy. The two off-grid rows report the share of stored transitions leaving the grid and the largest such excursion: these are bounded by the NLOpt constraint tolerance and the labor lower bound respectively, so they are float-sized rather than genuine extrapolation. The last two rows re-solve the same states against an interpolating cubic continuation instead of the linear one the solver uses, and report how far the optimal policy moves.